KnowledgePilot Test Document

A sample research brief for testing PDF text extraction, page tracking, and text chunking.

Document details

Document ID: KP-TEST-001
Version: 1.0
Publication date: 15 October 2026
Prepared for: KnowledgePilot development testing

Purpose

This document contains known facts distributed across three pages. Use it to check whether KnowledgePilot extracts text accurately and preserves page numbers. Later, the same document can be used to test chunking and retrieval.

Project overview

KnowledgePilot is a document research assistant. A user uploads a PDF, the application extracts its text, divides the text into smaller chunks, converts chunks into numerical embeddings, and retrieves relevant passages when the user asks a question. A language model can then use the retrieved passages to draft an answer grounded in the document.

Known fact for testing: The sample project name is KnowledgePilot, and its document ID is KP-TEST-001.

Page 2 — Document Processing

PDF text extraction

Text extraction converts the words stored in a PDF into text that a program can process. In this sample, the extraction test should find the phrase “page tracking helps preserve the origin of each passage.”

Chunking

Chunking divides long text into smaller passages. A chunk should be large enough to preserve a useful idea but small enough to retrieve without bringing in too much unrelated information. Overlap means that a small part of one chunk is repeated at the beginning of the next chunk. This can help preserve meaning when a sentence or idea crosses a chunk boundary.

Recommended test settings

Setting	Test value
Chunk size	120 words
Chunk overlap	20 words
Document pages	3
Document ID	KP-TEST-001

Known fact for testing: The recommended chunk overlap in this sample is 20 words.

Page 3 — Retrieval and Evaluation

Retrieval

Retrieval searches a collection of chunks and selects passages that are relevant to a user's question. A retrieval system may use keyword matching, semantic similarity, or a combination of both. Retrieval does not guarantee that the exact answer is present; it returns the passages it considers most relevant.

Answer grounding

When the retrieved passages contain enough evidence, the language model can use them as context to produce an answer. If the passages do not contain enough evidence, a reliable document assistant should say that the uploaded document does not provide sufficient information instead of inventing an answer.

Evaluation checklist

1. Confirm that all three pages are detected.
2. Confirm that the document ID KP-TEST-001 is extracted.
3. Confirm that the chunk overlap value is 20 words.
4. Confirm that page references remain attached to extracted passages.
5. Confirm that unsupported questions are not answered with invented document facts.

Known fact for testing: The sample document has three pages.